

FILTERING

(DIGITAL IMAGE PROCESSING)

Faculty of Information Technology
Ton Duc Thang University

August 2023

Image transformations

- As with any function, we can apply operators to an image



$$g(x,y) = f(x,y) + 20$$




$$g(x,y) = f(-x,y)$$

- We'll talk about a special kind of operator, convolution (linear filtering)

Question: Noise reduction

- Given a camera and a still scene, how can you reduce noise?



Source: S. Seitz

Take lots of images and
average them!

What's the next
best thing?

Image filtering

- Modify the pixels in an image based on some function of a **local neighborhood** of each pixel

10	5	3
4	5	1
1	1	7

Local image data

Some
function



	7	

Modified image data

Source: L. Zhang

Image filtering

- **Filtering:**

- Form a new image whose pixels are a combination original pixel values

Goals:

- Extract useful information from the images
 - Features (edges, corners, blobs...)
- Modify or enhance image properties:
 - super-resolution; in-painting; de-noising

De-noising



Salt and pepper noise

Super-resolution




In-painting



Bertamio et al

Linear filtering

- One simple version: linear filtering (cross-correlation, convolution)
 - Replace each pixel by a **linear combination of its neighbors**
- The prescription for the linear combination is called the “**kernel**” (or “mask”, “filter”)



Source: L. Zhang

Correlation of Discrete Signals

Correlation is a measure of how similar signals are

$$\text{corr}_{x,y} = \sum_{n=-\infty}^{\infty} x[n] y[n]$$

$$\text{corr}_{x,y} = \sum_{n=0}^{N-1} x[n] y[n]$$

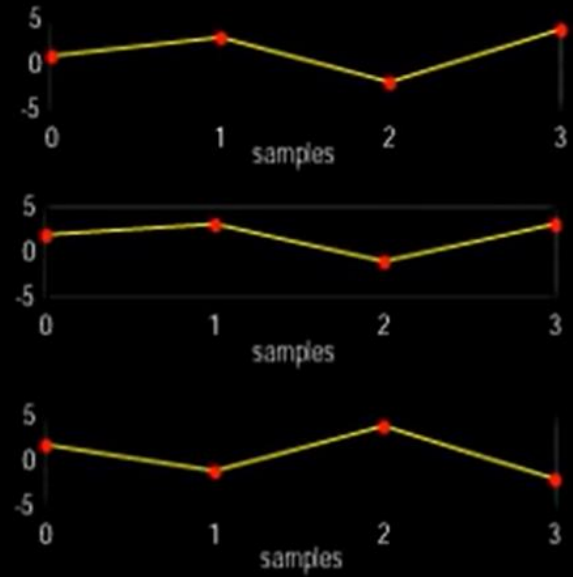
$$x = [1 \ 3 \ -2 \ 4]$$

$$\begin{aligned} \text{corr}_{x,y} &= x[0]y[0] + x[1]y[1] + x[2]y[2] + x[3]y[3] \\ &= (1)(2) + (3)(3) + (-2)(-1) + (4)(3) \\ &= 2 + 9 + 2 + 12 = 25 \end{aligned}$$

$$y = [2 \ 3 \ -1 \ 3]$$

$$\begin{aligned} \text{corr}_{y,z} &= y[0]z[0] + y[1]z[1] + y[2]z[2] + y[3]z[3] \\ &= 2(2) + (3)(-1) + (-1)(4) + (3)(-2) \\ &= 4 - 3 - 4 - 6 = -9 \end{aligned}$$

$$z = [2 \ -1 \ 4 \ -2]$$



Cross-correlation

- Let F be the image, H be the kernel (of size $2k + 1 \times 2k + 1$), and G be the output image:

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

This is called a **cross-correlation** operation:

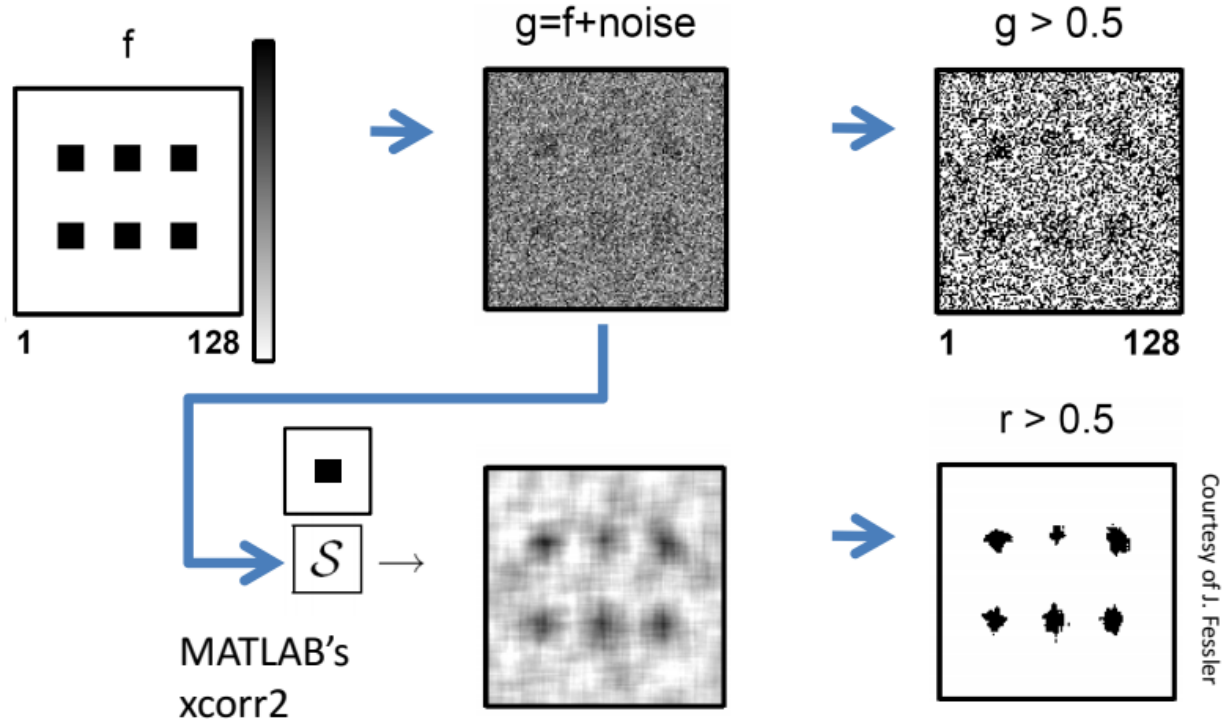
$$G = H \otimes F$$

- When the aperture is partially outside the image, the operation interpolates outlier pixel values according to the specified border mode (refers [\[1\]](#))

Cross-correlation (ct)

- How similar the kernel is to the image at any point [\[2\]](#)
 - Used for image alignment and simple image matching
- Refers [\[3\]](#) [\[4\]](#) more about template matching and normalized cross-correlation.

(Cross) correlation – example



Courtesy of J. Fessler



Cross Correlation Application: Vision system for TV remote control

- uses template matching







Figure from "Computer Vision for Interactive Computer Graphics," W.Freeman et al, IEEE Computer Graphics and Applications, 1998 copyright 1998, IEEE

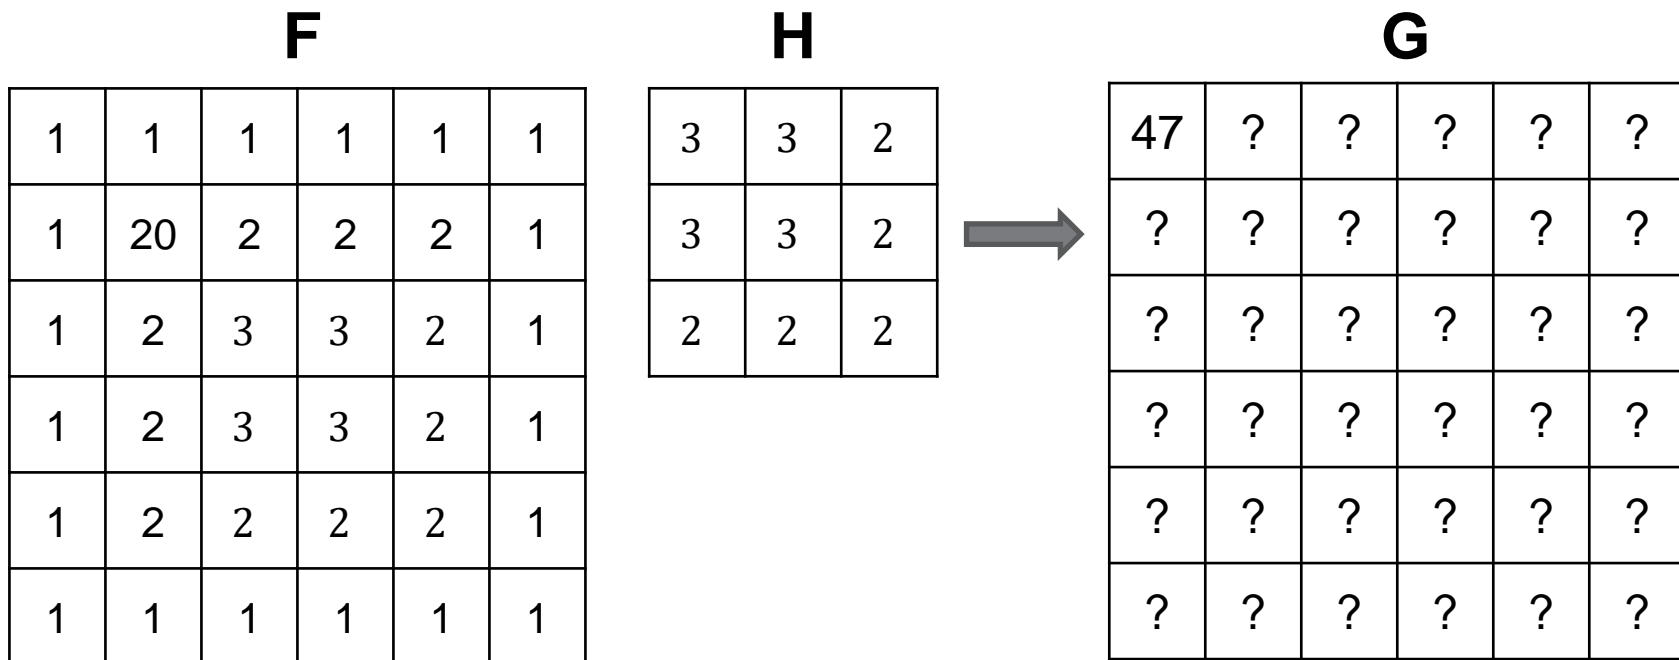
Ex. 1

- Apply cross-correlation operation into the following image F:



Ex. 2


- Apply **normalized** cross-correlation operation to locate the best matching of H in the image F (zero padding for the border pixels):



Convolution Operation

$$y(t) = h(t) * x(t) = \int_{-\infty}^{\infty} x(\tau) \cdot h(t-\tau) d\tau$$

↑
convolution op.



115 / **Signal & System**

<https://www.youtube.com/watch?app=desktop&v= HATc2zAhcY>

Convolution

- Same as cross-correlation, except that the kernel is “flipped” (horizontally and vertically)

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

This is called a **convolution** operation:

$$G = H * F$$

- Convolution is commutative and associative

2D Convolution

- $g(x,y) = h(x,y) * f(x,y)$

- f, g: input/output
- h: mask/filter/kernel



- **Flip** the mask (horizontally and vertically) only once
- Slide the mask onto the image.
- Multiply the corresponding elements and then add them
- Repeat this procedure until all values of the image has been calculated.

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0


Image

4		


Convolved
Feature

Example


- http://www.songho.ca/dsp/convolution/convolution2d_example.html



CONVOLUTION



Adapted from F. Durand



https://www.allaboutcircuits.com/uploads/articles/Fig2_2D_Conv.jpg


Convolution applications

- Blur image
- Remove noise
- Sharpening
- Smoothing
- Edge detection
- ...

<https://www.geeksforgeeks.org/python-opencv-filter2d-function/>

Ex. 3

- Apply convolution operation into the following image F:



Convolution vs. (Cross) Correlation

- A **convolution** is an integral that expresses the amount of overlap of one function as it is shifted over another function.
 - convolution is a filtering operation
- **Correlation** compares the *similarity of two sets of data*. Correlation computes a measure of similarity of two input signals as they are shifted by one another. The correlation result reaches a maximum at the time when the two signals match best .
 - correlation is a measure of relatedness of two signals

Linear filters: examples



Original



0	0	0
0	1	0
0	0	0



Identical image

Source: D. Lowe

Linear filters: examples



Original




0	0	0
1	0	0
0	0	0



Shifted left By 1 pixel

Source: D. Lowe

Linear filters: examples



Original



$\frac{1}{9}$


1	1	1
1	1	1
1	1	1



Blur (with a mean filter)

Source: D. Lowe

Mean filtering



Ex. 4

- Apply the filtering (cross-correlation) into the following image F (zero padding at the borders):

F

1	1	1	1	1	1
1	180	180	180	180	1
1	180	96	96	180	1
1	180	96	96	180	1
1	180	180	180	180	1
1	1	1	1	1	1

H

0	0	0
1	0	0
0	0	0



G

?	?	?	?	?	?
?	?	?	?	?	?
?	?	?	?	?	?
?	?	?	?	?	?
?	?	?	?	?	?
?	?	?	?	?	?

Ex. 5

- Apply the filtering (cross-correlation) into the following image F (zero padding at the borders):

F

1	1	1	1	1	1
1	180	180	180	180	1
1	180	96	96	180	1
1	180	96	96	180	1
1	180	180	180	180	1
1	1	1	1	1	1

H


0	0	0
0	0	1
0	0	0




G

?	?	?	?	?	?
?	?	?	?	?	?
?	?	?	?	?	?
?	?	?	?	?	?
?	?	?	?	?	?
?	?	?	?	?	?

Linear filters: examples




$$* \left(\begin{bmatrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} \right) =$$


Original


Sharpening
filter
(accentuates
edges)

Source: D. Lowe

SHARPENING



before



after

Source: D. Lowe

Sharpening

- emphasizes differences in adjacent pixel values
- accentuating the edges of the image
- add contrast to edges

Sharpen convolution

Input image:

0	128	255
64	192	0
255	0	64

Kernel:

0	-1	0
-1	5	-1
0	-1	0

anchor

Output image:

0	193	255
0	255	0
255	0	255

(assuming transparent border)

<https://i.stack.imgur.com/XXBUN.png>

Ex. 6

- Apply the filtering (cross-correlation) into the following image F (zero padding at the borders):

F

20	20	20	20	20	20
20	120	120	120	120	20
20	120	20	20	120	20
20	120	20	20	120	20
20	120	120	120	120	20
20	20	20	20	20	20

H

0	-1	0
-1	5	-1
0	-1	0




G

?	?	?	?	?	?
?	?	?	?	?	?
?	?	?	?	?	?
?	?	?	?	?	?
?	?	?	?	?	?
?	?	?	?	?	?



SMOOTHING WITH BOX FILTER

REVISITED



Source: D. Forsyth

GAUSSIAN KERNEL



$\frac{1}{16} \times$

1	2	1
2	4	2
1	2	1

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

<https://theailearner.com/2019/05/06/gaussian-blurring/>

Source: C. Rasmussen

Discrete approximation of the Gaussian kernels

1/16

1	2	1
2	4	2
1	2	1

1/273

1	4	7	4	1
4	16	26	16	4
7	26	41	26	7
4	16	26	16	4
1	4	7	4	1

1/1003

0	0	1	2	1	0	0
0	3	13	22	13	3	0
1	13	59	97	59	13	1
2	22	97	159	97	22	2
1	13	59	97	59	13	1
0	3	13	22	13	3	0
0	0	1	2	1	0	0

https://www.researchgate.net/figure/Discrete-approximation-of-the-Gaussian-kernels-3x3-5x5-7x7_fig2_325768087

Gaussian blur

- Use a weighted mean: the values near the center pixel will have a higher weight
 - probably get a less blurred image but a natural blurred image because it handles the edge values very well


<https://theailearner.com/tag/gaussian-filter/>

Gaussian kernel - properties

- Gaussian kernel is linearly separable: can break any 2-d filter into two 1-d filters
 - Applying multiple successive Gaussian kernels is equivalent to applying a single, larger Gaussian blur

$$\frac{1}{16} \begin{pmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{pmatrix} = \frac{1}{16} \begin{pmatrix} 1 \\ 2 \\ 1 \end{pmatrix} \begin{pmatrix} 1 & 2 & 1 \end{pmatrix}$$

- Gaussian kernel weights(1-D) can be obtained quickly using the Pascal's Triangle




<https://theailearner.com/tag/gaussian-filter/>


Gaussian blur – code

```
1 import cv2
2 img = cv2.imread('D:/downloads/opencv_logo.PNG')
3
4 # Creates a 1-D Gaussian kernel
5 a = cv2.getGaussianKernel(5,1)
6
7 # Apply the above Gaussian kernel. Here, I
8 # have used the same kernel for both X and Y
9 b = cv2.sepFilter2D(img,-1,a,a)
10
11 # Display the Image
12 cv2.imshow('a',b)
13 cv2.waitKey(0)
```


Gaussian filters



$\sigma = 1$
pixel



$\sigma = 5$
pixels




$\sigma = 10$
pixels



$\sigma = 30$
pixels

Sharpening revisited



Source: S. Lazebnik

Let's add it back:



$$F + \alpha(F - H * F)$$

image blurred image

Sharpen filter



Convolution in the real world

Camera shake



Source: Fergus, *et al.* "Removing Camera Shake from a Single Photograph", SIGGRAPH 2006

Bokeh: Blur out-of-focus regions of an image.

Source: <http://lullaby.homepage.dk/diy-camera/bokeh.html>



Rank filters

- Rank filters assign the **k-th** value of the gray levels from the window consisting of M pixels sorted in ascending order [code]
 - The special cases $k = 1$, $k = M$ (MIN and MAX filter) : erosion and dilation
 - $k = (M + 1)/2$: median filter
- Generalisation of flat dilation/erosion: in lieu of min or max value in window, use the k -th ranked value
- Increases robustness against noise
 - Best-known example: median filter for noise reduction
- Concept useful for both gray-level and binary images
- All rank filters are commutative with thresholding

Rank filters - benefits

- image quality enhancement, e.g., image smoothing, sharpening
- image pre-processing, e.g., noise reduction, contrast enhancement
- feature extraction, e.g., border detection, isolated point detection
- image post-processing, e.g., small object removal, object grouping, contour smoothing

Median filter

- Gray-level median filter

$$g[x, y] = \text{median}\left[W\left\{f[x, y]\right\}\right] := \text{median}(f, W)$$

- Binary images: majority filter

$$g[x, y] = \text{MAJ}\left[W\left\{f[x, y]\right\}\right] := \text{majority}(f, W)$$

- Self-duality

$$\text{median}(f, W) = -\left[\text{median}(-f, W)\right]$$

$$\text{majority}(f, W) = \text{NOT}\left[\text{majority}(\text{NOT}[f], W)\right]$$

Median filter

Input Window:

50	10	20
30	70	90
40	60	80

Sorter Output:

Low → 10
→ 20
→ 30
→ 40
Median → 50
→ 60
→ 70
→ 80
High → 90

For a RO filter with
Order = 5 (median)

Output Pixel:


—	—	—
—	50	—
—	—	—

https://www.researchgate.net/figure/Graphic-Depiction-of-Rank-Order-Filter-Operation_fig6_268373873

Majority filter: example




Binary image with
5% 'Salt&Pepper' noise



3x3 majority filter



20% 'Salt&Pepper' noise



3x3 majority filter

Median filter: example



Original
image



5% 'Salt&Pepper'
noise

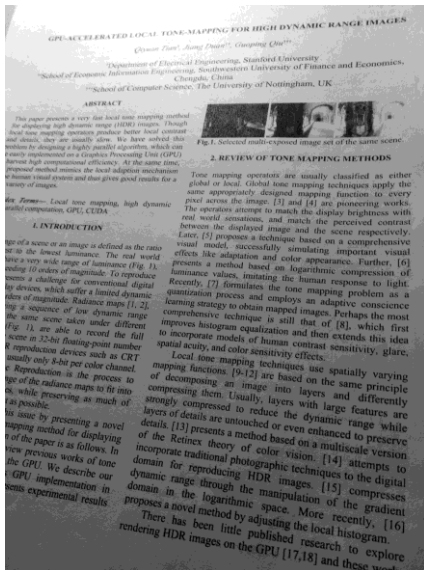


3x3 median
filtering



7x7 median
filtering

Example: non-uniform lighting compensation



Original image
1632x1216 pixels

Dilation (local max)
61x61 structuring element

Rank filter
10th brightest pixel
61x61 structuring element

References

1. https://docs.opencv.org/4.x/d4/d86/group_imgproc_filter.html
2. <https://www.youtube.com/watch?app=desktop&v=kGHz-cEyjiE>
3. https://docs.opencv.org/4.x/d4/dc6/tutorial_py_template_matching.html
4. <https://www.youtube.com/watch?app=desktop&v=kGHz-cEyjiE>
5. <https://vincmazet.github.io/bip/filtering/convolution.html>